

DSII Final Project

Yiru Gong, yg2832; Yiwen Zhao, yz4187; Jiaqi Chen, jc5681

2022-05-12

Contents

Data Input	2
Exploratory analysis	3
Data split	5
Model fitting	5
Penalized logistic regression (GLMNET)	5
GAM	7
LDA	9
Random Forest	9
Support Vector Machine	11
Neural Network	14
Model Comparison	16
CV Compare	16
Test data performance	18

```
data = read.csv('Covid19_vacc_predict_handout.csv')
data = data %>%
  na.omit() %>%
  dplyr::select(-id) %>%
  mutate(
    atlas_type_2015_mining_no = factor(atlas_type_2015_mining_no),
    covid_vaccination = factor(covid_vaccination),
    hum_region = factor(hum_region),
    sex_cd = factor(sex_cd),
    race_cd = factor(race_cd),
    lang_spoken_cd = factor(lang_spoken_cd),
    atlas_low_education_2015_update = factor(atlas_low_education_2015_update)
  )
dfSummary(data[,c(5,7,8,10,11,17,18)])
```

Data Frame Summary
Dimensions: 8308 x 7
Duplicates: 7802

No	Variable	Stats / Values	Freqs (% of Valid)	Graph	Valid	Missing
1	atlas_type_2015_mining1. [factor]	1. no 2. 1	8177 (98.4%) 131 (1.6%)	IIIIIIIIIIIIIIIIIIII I	8308 (100.0%)	0 (0.0%)
2	covid_vaccination [factor]	1. no_vacc 2. vacc	6682 (80.4%) 1626 (19.6%)	IIIIIIIIIIIIIIIIIIII III	8308 (100.0%)	0 (0.0%)

No	Variable	Stats / Values	Freqs (% of Valid)	Graph	Valid	Missing
3	hum_region [factor]	1. CALIFORNIA/NEVADA 2. CENTRAL 3. CENTRAL WEST 4. EAST 5. EAST CENTRAL 6. FLORIDA 7. GREAT LAKES/CENTRAL NORTH 8. GULF STATES 9. INTERMOUNTAIN 10. MID-ATLANTIC/NORTH CAROLI [5 others]	299 (3.6%) 551 (6.6%) 238 (2.9%) 491 (5.9%) 1370 (16.5%) 607 (7.3%) 1111 (13.4%) 454 (5.5%) 220 (2.6%) 845 (10.2%) 2122 (25.5%)	I	8308 (100.0%)	0 (0.0%)
4	sex_cd [factor]	1. F 2. M	4527 (54.5%) 3781 (45.5%)	IIIIIIII IIIIIIII	8308 (100.0%)	0 (0.0%)
5	lang_spoken_cd [factor]	1. * 2. CHI 3. CRE 4. ENG 5. KOR 6. OTH 7. SPA 8. VIE	10 (0.1%) 13 (0.2%) 4 (0.0%) 7957 (95.8%) 7 (0.1%) 34 (0.4%) 276 (3.3%) 7 (0.1%)		8308 (100.0%)	0 (0.0%)
6	atlas_low_education_2015 Update [factor]		7769 (93.5%) 539 (6.5%)	IIIIIIIIIIIIIIII I	8308 (100.0%)	0 (0.0%)
7	race_cd [factor]	1. 0 2. 1 3. 2 4. 3 5. 4 6. 5 7. 6	160 (1.9%) 7317 (88.1%) 558 (6.7%) 80 (1.0%) 56 (0.7%) 129 (1.6%) 8 (0.1%)	IIIIIIIIIIIIIIII I	8308 (100.0%)	0 (0.0%)

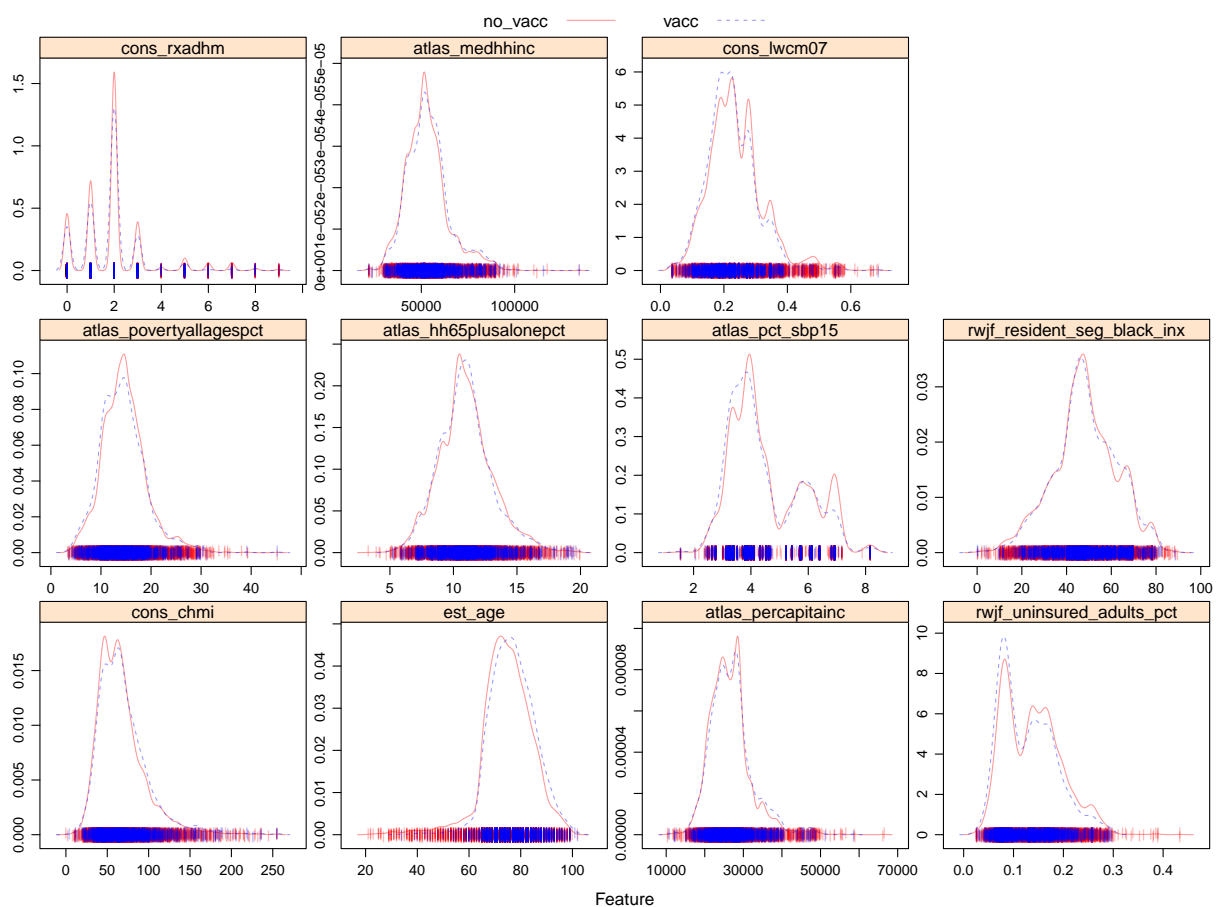
```
data2 = model.matrix(covid_vaccination ~ ., data)[ , -1]
```

Exploratory analysis


```
theme1 <- transparentTheme(trans = .4)
trellis.par.set(theme1)

#figure 1
featurePlot(x = data[, -c(5, 7, 8, 10, 11, 17, 18)],
            y = data$covid_vaccination,
            scales = list(x = list(relation = "free"),
                          y = list(relation = "free")),
```

```
plot = "density", pch = "|",
auto.key = list(columns = 2))
```



```
#correlation
corrplot(cor(data[, -c(5, 7, 8, 10, 11, 17, 18)]), method = "circle", type = "full")
```



Data split

```
set.seed(1)
rowTrain <- createDataPartition(y = data$covid_vaccination,
                                p = 0.7,
                                list = FALSE)

x = data2[rowTrain,]
y = data$covid_vaccination[rowTrain]
x2 = data2[-rowTrain,]
y2 = data$covid_vaccination[-rowTrain]

save(x,y,x2,y2,file = "split_data.Rdata")
```

Model fitting

Penalized logistic regression (GLMNET)

```
ctrl <- trainControl(method = "cv",
                     summaryFunction = twoClassSummary,
```

```

classProbs = TRUE)

glmnGrid <- expand.grid(.alpha = seq(0, 1, length = 21),
                      .lambda = exp(seq(-8, -1, length = 50)))

set.seed(1)
model.glmn <- train(x, y,
                    method = "glmnet",
                    tuneGrid = glmnGrid,
                    metric = "ROC",
                    trControl = ctrl)

model.glmn$bestTune

```

```

##      alpha      lambda
## 89  0.05 0.07642629


```

```

myCol<- rainbow(25)
myPar <- list(superpose.symbol = list(col = myCol),
              superpose.line = list(col = myCol))

plot(model.glmn, par.settings = myPar, xTrans = function(x) log(x))

```



GAM

```

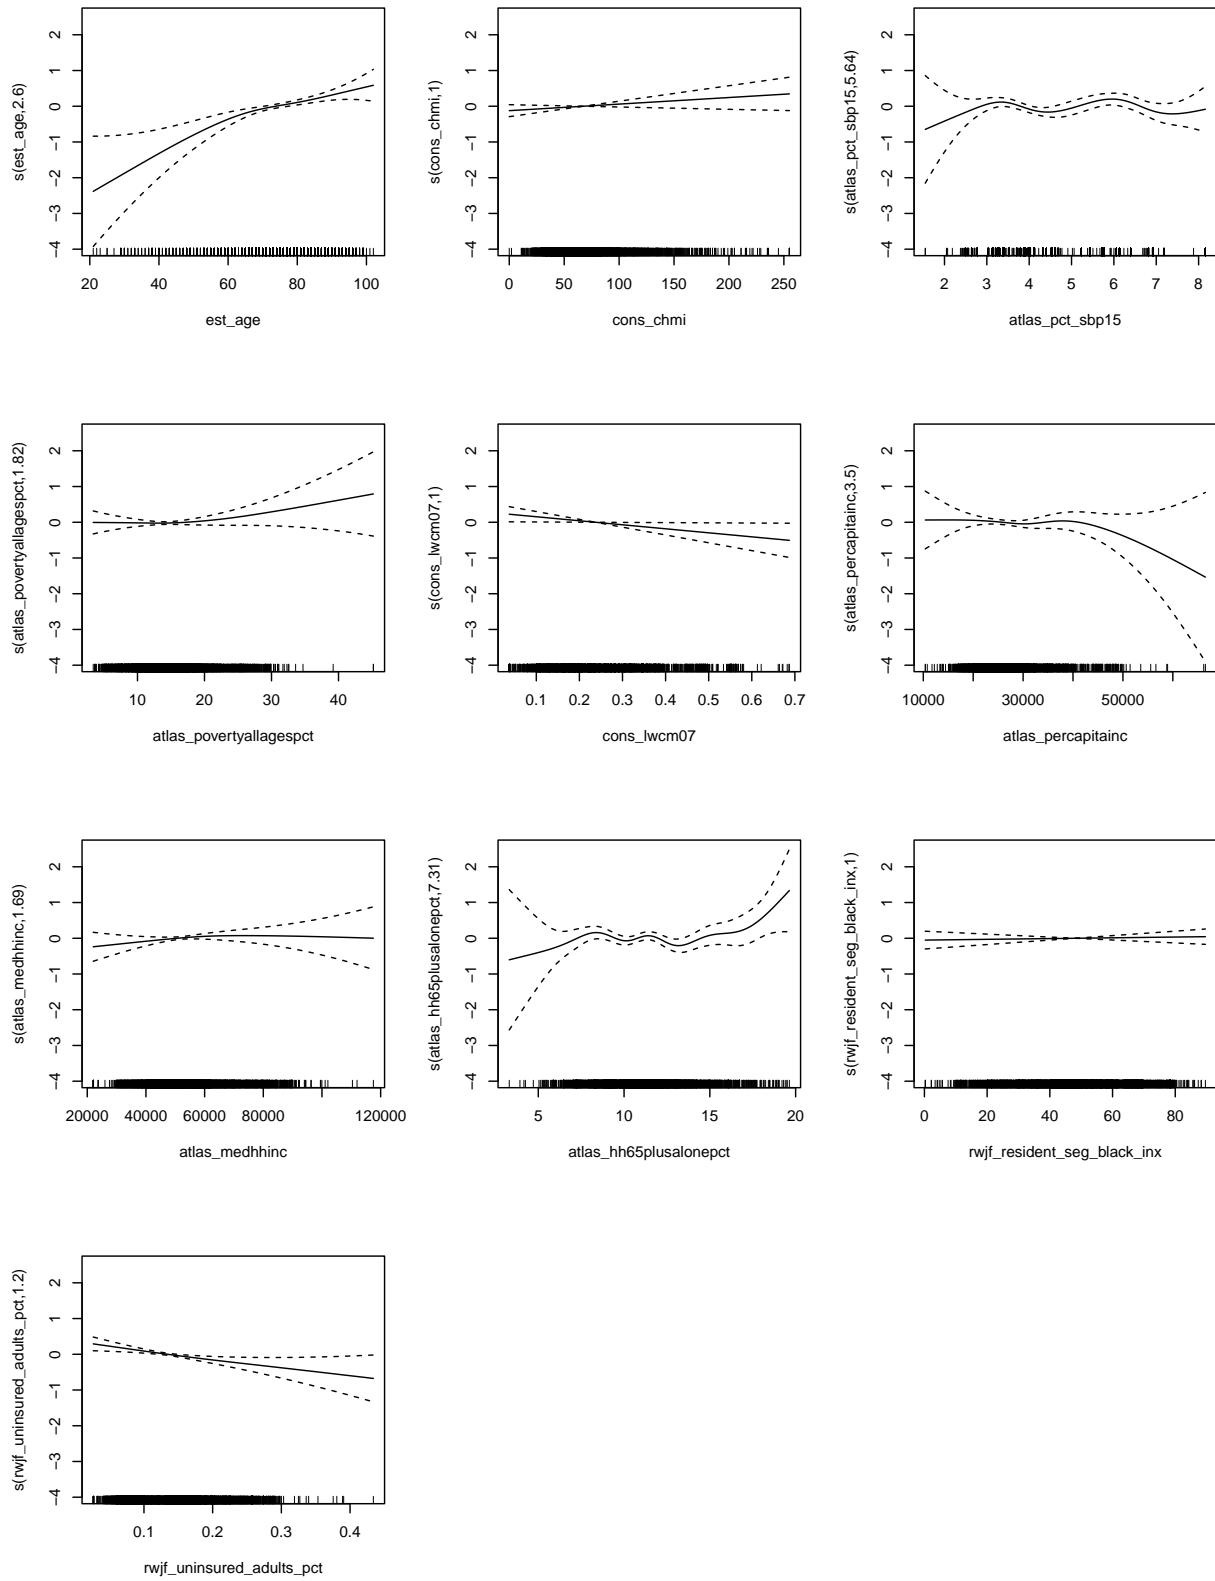
set.seed(1)
model.gam <- train(data[rowTrain,-c(7:8)], y,
                    method = "gam",
                    metric = "ROC",
                    trControl = ctrl)
### row 8: hum_region report error

model.gam$finalModel

##
## Family: binomial
## Link function: logit
##
## Formula:
## .outcome ~ sex_cd + atlas_low_education_2015_update + race_cd +
##   cons_rxadhm + s(est_age) + s(cons_chmi) + s(atlas_pct_sbp15) +
##   s(atlas_povertyallagespct) + s(cons_lwcm07) + s(atlas_percapitainc) +
##   s(atlas_medhhinc) + s(atlas_hh65plusalonepct) + s(rwjf_resident_seg_black_inx) +
##   s(rwjf_uninsured_adults_pct)
##
## Estimated degrees of freedom:
## 2.60 1.00 5.64 1.82 1.00 3.50 1.69
## 7.31 1.00 1.20 total = 36.76
##
## UBRE score: -0.02449249



# fig 2
par(mfrow=c(4,3))
plot(model.gam$finalModel)

```



LDA

```
lda.fit <- lda(y~x)
plot(lda.fit)
```



```
set.seed(1)
model.lda <- train(x, y,
  method = "lda",
  metric = "ROC",
  trControl = ctrl)
```


Random Forest

```
# rf.ctrl <- trainControl(method = "cv",
#                           classProbs = TRUE,
#                           summaryFunction = twoClassSummary)

rf.grid <- expand.grid(mtry = 1:8,
  splitrule = "gini",
  min.node.size = seq(from = 2, to = 10,
    by = 2))
```


```
set.seed(1)
rf.fit <- train(covid_vaccination ~ . ,
               data,
               subset = rowTrain,
               method = "ranger",
               tuneGrid = rf.grid,
               metric = "ROC",
               trControl = ctrl)

ggplot(rf.fit, highlight = TRUE)
```



```
# variable importance
set.seed(1)
rf2.final.per <- ranger(covid_vaccination ~ . ,
                       data[rowTrain,],
                       mtry = rf.fit$bestTune[[1]],
                       min.node.size = rf.fit$bestTune[[3]],
                       splitrule = "gini",
                       importance = "permutation",
                       scale.permutation.importance = TRUE)

par(mar = c(3,12,3,3))
barplot(sort(ranger::importance(rf2.final.per), decreasing = FALSE),
        las = 2, horiz = TRUE, cex.names = 0.7,
        col = colorRampPalette(colors = c("cyan","blue"))(17))
```



Support Vector Machine

```
data$covid_vaccination <- factor(data$covid_vaccination, c("vacc", "no_vacc"))
dat <- data[-c(5,8,10,11,17,18)]
summary(dat)
```


```
##      cons_chmi      est_age      atlas_percapitainc rwjf_uninsured_adults_pct
## Min.   : 0.00   Min.   : 21.00   Min.   :10399      Min.   :0.02616
## 1st Qu.: 47.00   1st Qu.: 70.00   1st Qu.:23056      1st Qu.:0.08593
## Median : 62.00   Median : 75.00   Median :26132      Median :0.13357
## Mean   : 67.21   Mean   : 75.18   Mean   :26685      Mean   :0.13645
## 3rd Qu.: 79.00   3rd Qu.: 81.00   3rd Qu.:28949      3rd Qu.:0.17323
## Max.   :255.00   Max.   :102.00   Max.   :66522      Max.   :0.43395
## atlas_povertyallagespct covid_vaccination atlas_hh65plusalonepct
## Min.   : 3.40             vacc   :1626      Min.   : 3.309
## 1st Qu.:11.60             no_vacc:6682      1st Qu.: 9.626
## Median :14.40                                     Median :10.878
## Mean   :14.61                                     Mean   :10.993
## 3rd Qu.:17.00                                     3rd Qu.:12.155
## Max.   :45.20                                     Max.   :19.960
## atlas_pct_sbp15 rwjf_resident_seg_black_inx cons_rxadhm atlas_medhhinc
## Min.   :1.546   Min.   : 0.2584      Min.   :0.000   Min.   : 22045
## 1st Qu.:3.525   1st Qu.:40.3734      1st Qu.:1.000   1st Qu.: 45813
```

```
## Median :4.110   Median :47.8798           Median :2.000   Median : 51865
## Mean    :4.559   Mean    :48.1327           Mean    :1.921   Mean    : 53259
## 3rd Qu.:5.710   3rd Qu.:57.8018           3rd Qu.:2.000   3rd Qu.: 58742
## Max.    :8.160   Max.    :89.6102           Max.    :9.000   Max.    :134609
## cons_lwcm07
## Min.    :0.03724
## 1st Qu.:0.18190
## Median :0.22509
## Mean    :0.23641
## 3rd Qu.:0.28204
## Max.    :0.68722
```

```
# SVM with Linear Kernal
# ctrl1 <- trainControl(method = "cv")
set.seed(1)
svml.fit <- train(covid_vaccination ~ . ,
                  data = dat[rowTrain,],
                  method = "svmLinear",
                  metric = "ROC",
                  # preProcess = c("center", "scale"),
                  tuneGrid = data.frame(C = exp(seq(-5,2,len=50))),
                  trControl = ctrl1)
```

```
## maximum number of iterations reached 0.003289937 -0.003225867maximum number of iterations reached 0.
```

```
plot(svml.fit, highlight = TRUE, xTrans = log)
```



```
svml.fit$bestTune
```

```
##          C
## 47 4.81352
```

```
# SVM with Radial Kernel
```


```
svmr.grid <- expand.grid(C = exp(seq(-1,4,len=20)),
                        sigma = exp(seq(-6,-2,len=20)))
```

```
set.seed(1)
svmr.fit <- train(covid_vaccination ~ ., dat,
                 subset = rowTrain,
                 method = "svmRadialSigma",
                 metric = "ROC",
                 tuneGrid = svmr.grid,
                 trControl = ctrl)
```

```
## maximum number of iterations reached 0.001902067 -0.001852983maximum number of iterations reached 0.
```

```
myCol <- rainbow(20)
myPar <- list(superpose.symbol = list(col = myCol),
             superpose.line = list(col = myCol))

plot(svmr.fit, highlight = TRUE, par.settings = myPar)
```



```
svmr.fit$bestTune
```

```
##          sigma          C
## 392 0.025117 54.59815
```

Neural Network

```
## tuning
set.seed(1)
runs <- tuning_run("keras_grid_search.R",
                    flags = list(
                      nodes_layer1 = c(64, 128, 256),
                      nodes_layer2 = c(64, 128, 256),
                      nodes_layer3 = c(64, 128, 256),
                      dropout_layer1 = c(0.2, 0.3, 0.4),
                      dropout_layer2 = c(0.2, 0.3, 0.4),
                      dropout_layer3 = c(0.2, 0.3, 0.4)),
                    confirm = FALSE,
                    echo = FALSE,
                    sample = 0.01) # try more after class

best = runs[which.max(runs$metric_val_accuracy),]
best


## $ run_dir          : chr "runs/2022-05-12T15-14-35Z"
## $ metric_loss      : num 0.525
## $ metric_accuracy  : num 0.793
## $ metric_val_loss  : num 0.495
## $ metric_val_accuracy: num 0.808
## $ metric_lr        : num 0.001
## $ flag_nodes_layer1 : int 64
## $ flag_nodes_layer2 : int 64
## $ flag_nodes_layer3 : int 128
## $ flag_dropout_layer1: num 0.4
## $ flag_dropout_layer2: num 0.2
## $ flag_dropout_layer3: num 0.3
## $ epochs           : int 30
## $ epochs_completed  : int 16
## $ metrics           : chr "(metrics data frame)"
## $ model              : chr "(model summary)"
## $ loss_function      : chr "categorical_crossentropy"
## $ optimizer          : chr "<keras.optimizer_v2.rmsprop.RMSprop>"
## $ learning_rate     : num 0.001
## $ script             : chr "keras_grid_search.R"
## $ start              : POSIXct[1:1], format: "2022-05-12 15:14:35"
## $ end                : POSIXct[1:1], format: "2022-05-12 15:14:41"
## $ completed         : logi TRUE
## $ output             : chr "(script ouptut)"
## $ source_code        : chr "(source archive)"
## $ context            : chr "local"
## $ type               : chr "training"
```

```

y_c = ifelse(y=="vacc",1,0)
y_c <- to_categorical(y_c, 2)
y2_c = ifelse(y2=="vacc",1,0)
y2_c <- to_categorical(y2_c, 2)

model.nn <- keras_model_sequential() %>%
  layer_dense(units = best$flag_nodes_layer1, activation = "relu", input_shape = ncol(x)) %>%
  layer_batch_normalization() %>%
  layer_dropout(rate = best$flag_dropout_layer1) %>%
  layer_dense(units = best$flag_nodes_layer2, activation = "relu") %>%
  layer_batch_normalization() %>%
  layer_dropout(rate = best$flag_dropout_layer2) %>%
  layer_dense(units = best$flag_nodes_layer3, activation = "relu") %>%
  layer_batch_normalization() %>%
  layer_dropout(rate = best$flag_dropout_layer3) %>%
  layer_dense(units = 2, activation = "sigmoid") %>%
  compile(loss = "categorical_crossentropy",
          optimizer = optimizer_rmsprop(),
          metrics = "accuracy")
fit.nn = model.nn %>%
  fit(x = x,
      y = y_c,
      epochs = 30,
      batch_size = 256,
      validation_split = 0.2,
      callbacks = list(callback_early_stopping(patience = 5),
                       callback_reduce_lr_on_plateau()),
      verbose = 2)
plot(fit.nn)

```



```
## testing and evaluation
score <- model.nn %>% evaluate(x2, y2_c)
score
```

```
##      loss  accuracy
## 0.5046750 0.8044962
```

Model Comparison

CV Compare

```
res <- resamples(list(GLMNET = model.glmn,
                     GAM = model.gam,
                     LDA = model.lda,
                     RF = rf.fit,
                     SVML = svm1.fit,
                     SVMR = svmr.fit))
```


```
#KNN
summary(res)
```

```
##
```



```
## Call:
## summary.resamples(object = res)
##
## Models: GLMNET, GAM, LDA, RF, SVML, SVMR
## Number of resamples: 10
##
## ROC
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## GLMNET 0.5429205 0.5624063 0.5865405 0.5872412 0.6059043 0.6387389    0
## GAM    0.5400922 0.5726839 0.5844861 0.5816581 0.5903860 0.6107550    0
## LDA    0.5351253 0.5645383 0.5778804 0.5822242 0.6033940 0.6224134    0
## RF     0.5602789 0.5649677 0.5938578 0.5870003 0.6036325 0.6126106    0
## SVML   0.4205278 0.5348807 0.5615171 0.5498035 0.5852967 0.5957602    0
## SVMR   0.5078348 0.5310246 0.5599415 0.5526272 0.5734883 0.5912506    0
##
## Sens
##           Min. 1st Qu. Median     Mean 3rd Qu. Max. NA's
## GLMNET 1.0000000      1      1 1.0000000      1      1      0
## GAM    0.9978632      1      1 0.9995726      1      1      0
## LDA    1.0000000      1      1 1.0000000      1      1      0
## RF     1.0000000      1      1 1.0000000      1      1      0
## SVML   0.0000000      0      0 0.0000000      0      0      0
## SVMR   0.0000000      0      0 0.0000000      0      0      0
##
## Spec
##           Min. 1st Qu. Median     Mean 3rd Qu.     Max. NA's
## GLMNET      0      0      0 0.000000000      0 0.000000000      0
## GAM         0      0      0 0.000877193      0 0.000877193      0
## LDA         0      0      0 0.000000000      0 0.000000000      0
## RF          0      0      0 0.000000000      0 0.000000000      0
## SVML        1      1      1 1.000000000      1 1.000000000      0
## SVMR        1      1      1 1.000000000      1 1.000000000      0

# figure 4
bwplot(res, metric = "ROC")
```



Test data performance

```
# raw pred
glmnet.pred <- predict(model.glmnet, newdata = x2, type = "raw")
gam.pred <- predict(model.gam, newdata = data[-rowTrain,-c(7:8)], type = "raw")
lda.pred <- predict(model.lda, newdata = x2, type = "raw")
rf.pred <- predict(rf.fit, newdata = data[-rowTrain,], type = "raw")

svml.pred <- predict(svml.fit, newdata = dat[-rowTrain,], type = "raw")
svmr.pred <- predict(svmr.fit, newdata = dat[-rowTrain,], type = "raw")

pred_test <- model.nn %>% predict(x2) %>% k_argmax() %>% as.matrix() %>% as.numeric()
nn.pred = ifelse(pred_test==0,"no_vacc","vacc")
nn.pred = factor(nn.pred,levels = c("no_vacc","vacc"))

# Confusion Matrix
cm.glmnet = confusionMatrix(data = glmnet.pred, reference = y2, positive = "vacc")$overall
cm.gam = confusionMatrix(data = gam.pred, reference = y2, positive = "vacc")$overall
cm.lda = confusionMatrix(data = lda.pred, reference = y2, positive = "vacc")$overall
cm.rf = confusionMatrix(data = rf.pred, reference = y2, positive = "vacc")$overall
cm.svml = confusionMatrix(data = svml.pred, reference = y2, positive = "vacc")$overall
cm.svmr = confusionMatrix(data = svmr.pred, reference = y2, positive = "vacc")$overall
cm.nn = confusionMatrix(data = nn.pred, reference = y2, positive = "vacc")$overall
```

```
cm_df = data.frame(GLMN = cm.glmn, GAM = cm.gam, LDA = cm.lda, RF = cm.rf, SVML = cm.svm1, SVMR = cm.svmr,
  knitr::kable(cm_df, digits = 4)
```

```
glmn.pred <- predict(model.glmn, newdata = x2, type = "prob")[,2]
gam.pred <- predict(model.gam, newdata = data[-rowTrain,-c(7:8)], type = "prob")[,2]
lda.pred <- predict(model.lda, newdata = x2, type = "prob")[,2]
rf.pred <- predict(rf.fit, newdata = data[-rowTrain,], type = "prob")[,2]
svml.pred <- predict(svml.fit, newdata = dat[-rowTrain,], type = "prob")[,2]
svmr.pred <- predict(svmr.fit, newdata = dat[-rowTrain,], type = "prob")[,2]
```

```
pred_test <- model.nn %>% predict(x2)
nn.pred = pred_test[,2]
```

```
roc.glmn <- roc(y2, glmn.pred)
roc.gam <- roc(y2, gam.pred)
roc.lda <- roc(y2, lda.pred)
roc.rf <- roc(y2, rf.pred)
roc.svm1 <- roc(y2,svml.pred)
roc.svmr <- roc(y2,svmr.pred)
roc.nn = roc(y2,nn.pred)
```

```
auc <- c(roc.glmn$auc[1],
  roc.gam$auc[1],
  roc.lda$auc[1],
  roc.rf$auc[1],
  roc.svm1$auc[1],
  roc.svmr$auc[1],
  roc.nn$auc[1])
```

```
modelNames <- c("glmn","gam","lda","rf","svml","svmr","nn")
```

```
# fig 5
ggroc(list(roc.glmn, roc.gam, roc.lda, roc.rf,roc.svm1, roc.svmr, roc.nn),
  legacy.axes = TRUE) +
  scale_color_discrete(labels = paste0(modelNames, " (", round(auc,3),")"),
    name = "Models (AUC)") +
  geom_abline(intercept = 0, slope = 1, color = "grey")
```

